

A decorative graphic on the left side of the slide, consisting of a series of vertical and diagonal lines of varying lengths, some ending in small circles, resembling a circuit board or a stylized tree structure.

SYNCHRONIZATION AND ATOMIC OPERATIONS

SYNCHRONIZATION

Race Conditions



- Race conditions arise when 2+ threads attempt to access the same memory location concurrently and at least one access is a write.

```
// race.cu
__global__ void race(int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    *x = i;
}

// main.cpp
int x;
race<<<1,128>>>(d_x);
cudaMemcpy(&x, d_x, sizeof(int), cudaMemcpyDeviceToHost);
```

SYNCHRONIZATION

Race Conditions



- Programs with race conditions may produce unexpected, seemingly arbitrary results
 - * Updates may be missed, and updates may be lost

```
// race.cu
__global__ void race(int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    *x = *x + 1;
}

// main.cpp
int x;
race<<<1,128>>>(d_x);
cudaMemcpy(&x, d_x, sizeof(int), cudaMemcpyDeviceToHost);
```

SYNCHRONIZATION

Synchronization



- **Accesses to shared locations need to be correctly synchronized (coordinated) to avoid race conditions**
- **In many common shared memory multithreaded programming models, one uses coordination objects such as locks to synchronize accesses to shared data**
- **CUDA provides several scalable synchronization mechanisms, such as efficient barriers and atomic memory operations.**
- **In general, always most efficient to design algorithms to avoid synchronization whenever possible.**

SYNCHRONIZATION

Synchronization



- Assume thread T₁ reads a value defined by thread T₀

```
// update.cu
__global__ void update_race(int* x, int* y)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i == 0) *x = 1;
    if (i == 1) *y = *x;
}

// main.cpp
update_race<<<1,2>>>>(d_x, d_y);
cudaMemcpy(&y, d_y, sizeof(int), cudaMemcpyDeviceToHost);
```

- Program needs to ensure that thread T₁ reads location after thread T₀ has written location.

SYNCHRONIZATION



Synchronization within Block

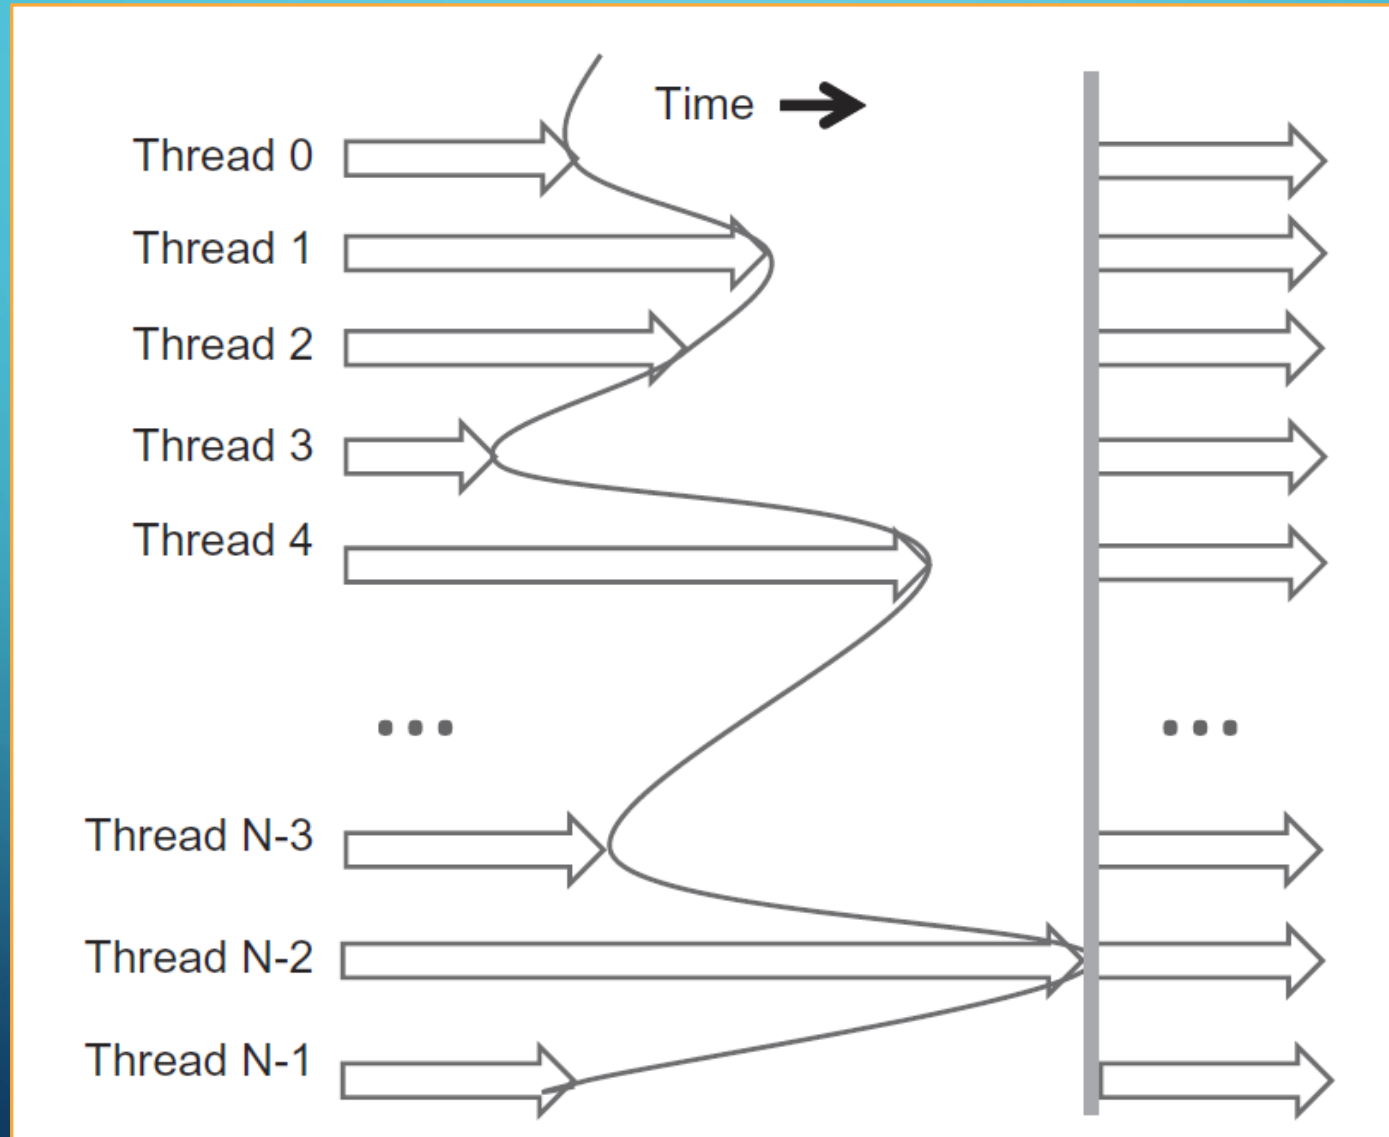
- Threads in same block: can use `__syncthreads()` to specify synchronization point that orders accesses

```
// update.cu
__global__ void update(int* x, int* y)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i == 0) *x = 1;
    __syncthreads();
    if (i == 1) *y = *x;
}

// main.cpp
update<<<1,2>>>>(d_x, d_y);
cudaMemcpy(&y, d_y, sizeof(int), cudaMemcpyDeviceToHost);
```

- **Important:** all threads within the block must reach the `__syncthreads()` statement

SYNCHRONIZATION



ATOMIC OPERATIONS

Introduction to Atomics



- Atom memory operations (**atomic functions**) are used to solve all kinds of synchronization and coordination problems in parallel computer systems.
- General concept is to provide a mechanism for a thread to update a memory location such that the update appears to happen atomically (without interruption) with respect to other threads.
- This ensures that all atomic updates issued concurrently are performed (often in some unspecified order) and that all threads can observe all updates.

ATOMIC OPERATIONS

Atomic Functions



- Atomic functions perform read-modify-write operations on data residing in global and shared memory

```
//example of int atomicAdd(int* addr, int val)
__global__ void update(unsigned int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    int j = atomicAdd(x, 1);    // j = *x; *x = j + i;
}
```

```
// main.cpp
int x = 0;
cudaMemcpy(d_x, x, cudaMemcpyHostToDevice);
update<<<1,128>>>>;
cudaMemcpy(&x, d_x, cudaMemcpyHostToDevice);
```

- Atomic functions guarantee that only one thread may access a memory location while the operation completes

ATOMIC OPERATIONS

Atomic Functions



- Synopsis of atomic function **atomicOP(a,b)** is typically

```
t1 = *a;          // read
t2 = t1 OP b;     // modify
*a = t2;          // write
return t;
```

- The hardware ensures that all statements are executed atomically without interruption by any other atomic functions.
- The atomic function returns the initial value, not the final value, stored at the memory location.

ATOMIC OPERATIONS

Atomic Functions



- The name atomic is used because the update is performed atomically: it cannot be interrupted by other atomic updates.
- The order in which concurrent atomic updates are performed is not defined, and may appear arbitrary.
- However, none of the atomic updates will be lost.
- Many different kinds of atomic operations
 - ✱ Add (add), Sub (subtract), Inc (increment), Dec (decrement)
 - ✱ And (bit-wise and), Or (bit-wise or) , Xor (bit-wise exclusive or)
 - ✱ Exch (Exchange)
 - ✱ Min (Minimum), Max (Maximum)
 - ✱ Compare-and-Swap

ATOMIC OPERATIONS

Histogram Example



```
// Compute histogram of colors in an image
//
// color - pointer to picture color data
// bucket - pointer to histogram buckets, one per color
//

__global__ void histogram(int n, int* color, int* bucket)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i < n)
    {
        int c = colors[i];
        atomicAdd(&bucket[c], 1);
    }
}
```

ATOMIC OPERATIONS

Performance Notes



- **Atomics are slower than normal accesses (loads, stores)**
- **Performance can degrade when many threads attempt to perform atomic operations on a small number of locations**
- **Possible to have all threads on the machine stalled, waiting to perform atomic operations on a single memory location.**

REFERENCES

- <https://developer.nvidia.com/cuda-education>