



# GPU Teaching Kit

Accelerated Computing



## Module 8.1 – Parallel Computation Patterns (Stencil)

### Convolution

# Objective

- To learn convolution, an important method
  - Widely used in audio, image and video processing
  - Foundational to stencil computation used in many science and engineering applications
  - Basic 1D and 2D convolution kernels

# Convolution as a Filter

- Often performed as a filter that transforms signal or pixel values into more desirable values.
  - Some filters smooth out the signal values so that one can see the big-picture trend
  - Others like Gaussian filters can be used to sharpen boundaries and edges of objects in images..

# Gaussian Blur

$$\frac{1}{273}$$

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1



# Convolution as a Edge Detection



-1	0	1
-2	0	2
-1	0	1




0	-1	2
-1	0	1
-2	-1	0



# Convolution – a computational definition

- An array operation where each output data element is a weighted sum of a collection of neighboring input elements
- The weights used in the weighted sum calculation are defined by an input mask array, commonly referred to as the *convolution kernel*
  - We will refer to these mask arrays as convolution masks to avoid confusion.
  - The value pattern of the mask array elements defines the type of filtering done
  - Our image blur example in Module 3 is a special case where all mask elements are of the same value and hard coded into the source code.


# 1D Convolution Example



- Commonly used for audio processing
  - Mask size is usually an odd number of elements for symmetry (5 in this example)
- The figure shows calculation of  $P[2]$


$$P[2] = N[0]*M[0] + N[1]*M[1] + N[2]*M[2] + N[3]*M[3] + N[4]*M[4]$$

# Calculation of P[3]





# Convolution Boundary Condition



- Calculation of output elements near the boundaries (beginning and end) of the array need to deal with “ghost” elements
  - Different policies (0, replicates of boundary values, etc.)

# A 1D Convolution Kernel with Boundary Condition Handling

- This kernel forces all elements outside the valid input range to 0

```
__global__ void convolution_1D_basic_kernel(float *N, float *M,  
                                           float *P, int Mask_Width, int Width)  
{  
    int i = blockIdx.x*blockDim.x + threadIdx.x;  
  
    float Pvalue = 0;  
    int N_start_point = i - (Mask_Width/2);  
  
    for (int j = 0; j < Mask_Width; j++) {  
        if (N_start_point + j >= 0 && N_start_point + j < Width) {  
            Pvalue += N[N_start_point + j]*M[j];  
        }  
    }  
  
    P[i] = Pvalue;  
}
```

# A 1D Convolution Kernel with Boundary Condition Handling

- This kernel forces all elements outside the valid input range to 0

```
__global__ void convolution_1D_basic_kernel(float *N, float *M,  
                                           float *P, int Mask_Width, int Width)  
{  
    int i = blockIdx.x*blockDim.x + threadIdx.x;  
  
    float Pvalue = 0;  
    int N_start_point = i - (Mask_Width/2);  
  
    if (i < Width) {  
        for (int j = 0; j < Mask_Width; j++) {  
            if (N_start_point + j >= 0 && N_start_point + j < Width) {  
                Pvalue += N[N_start_point + j]*M[j];  
            }  
        }  
  
        P[i] = Pvalue;  
    }  
}
```

# 2D Convolution

**N**

1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

**P**


1	2	3	4	5			
2	3	4	5	6			
3	4	321	6	7			
4	5	6	7	8			
5	6	7	8	5			

**M**

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1

1	4	9	8	5
4	9	16	15	12
4	16	25	24	21
8	15	24	21	16
5	12	21	16	5

# 2D Convolution – Ghost Cells



\_\_global\_\_

```
void convolution_2D_basic_kernel(unsigned char * in, unsigned char * mask, unsigned char * out,  
    int maskwidth, int w, int h) {
```

```
    int Col = blockIdx.x * blockDim.x + threadIdx.x;  
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (Col < w && Row < h) {
```

```
        int pixVal = 0;
```

```
        N_start_col = Col - (maskwidth/2);
```

```
        N_start_row = Row - (maskwidth/2);
```

```
        // Get the of the surrounding box
```

```
        for(int j = 0; j < maskwidth; ++j) {
```

```
            for(int k = 0; k < maskwidth; ++k) {
```

```
                int curRow = N_Start_row + j;
```

```
                int curCol = N_start_col + k;
```

```
                // Verify we have a valid image pixel
```

```
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w)  
                    pixVal += in[curRow * w + curCol] * mask[j*maskwidth+k];
```

```
            }
```


```
        }
```

```
    }
```

```
    // Write our new pixel value out
```

```
    out[Row * w + Col] = (unsigned char)(pixVal);
```

```
}
```



\_\_global\_\_

```
void convolution_2D_basic_kernel(unsigned char * in, unsigned char * mask, unsigned char * out,  
    int maskwidth, int w, int h) {  
    int Col = blockIdx.x * blockDim.x + threadIdx.x;  
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (Col < w && Row < h) {
```

```
        int pixVal = 0;
```

```
        N_start_col = Col - (maskwidth/2);  
        N_start_row = Row - (maskwidth/2);
```

N\_start\_row

```
        // Get the of the surrounding box
```

```
        for(int j = 0; j < maskwidth; ++j) {
```

```
            for(int k = 0; k < maskwidth; ++k) {
```

```
                int curRow = N_Start_row + j;
```

```
                int curCol = N_start_col + k;
```

```
                // Verify we have a valid image pixel
```

```
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w)  
                    pixVal += in[curRow * w + curCol] * mask[j*maskwidth+k];
```

```
            }
```

```
        }
```

```
    }
```

```
    // Write our new pixel value out
```

```
    out[Row * w + Col] = (unsigned char)(pixVal);
```

N\_start\_col

**N**

1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

**M**

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1

1	4	9	8	5
4	9	16	15	12
4	16	25	24	21
8	15	24	21	16
5	12	21	16	5

\_\_global\_\_

```
void convolution_2D_basic_kernel(unsigned char * in, unsigned char * mask, unsigned char * out,  
    int maskwidth, int w, int h) {  
    int Col = blockIdx.x * blockDim.x + threadIdx.x;  
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (Col < w && Row < h) {  
        int pixVal = 0;
```

```
        N_start_col = Col - (maskwidth/2);  
        N_start_row = Row - (maskwidth/2);
```

```
        // Get the of the surrounding box
```

```
        for(int j = 0; j < maskwidth; ++j) {  
            for(int k = 0; k < maskwidth; ++k) {  
  
                int curRow = N_Start_row + j;  
                int curCol = N_start_col + k;  
                // Verify we have a valid image pixel  
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w) {  
                    pixVal += in[curRow * w + curCol] * mask[j*maskwidth+k];  
                }  
            }  
        }  
    }
```

```
        // Write our new pixel value out
```

```
        out[Row * w + Col] = (unsigned char)(pixVal);  
    }
```



# CONSTANT MEMORY AND CACHING

- There are three interesting properties of the way the mask array  $M$  is used in convolution.
- First, the size of the  $M$  array is typically small. Most convolution masks are less than 10 elements in each dimension. Even in the case of a 3D convolution, the mask typically contains only less than 1000 elements.
- Second, the contents of  $M$  are not changed throughout the execution of the kernel.
- Third, all threads need to access the mask elements.
- Even better, all threads access the  $M$  elements in the same order, starting from  $M[0]$  and move by one element a time.
- These two properties make the mask array an excellent candidate for constant memory and caching.

# CONSTANT MEMORY AND CACHING

- Like global memory variables, constant memory variables are also visible to all thread blocks.
- The main difference is that a constant memory variable cannot be changed by threads during kernel execution.
- Furthermore, the size of the constant memory is quite small, currently at 64KB.
- To declare an M array in constant memory, the host code declares it as a global variable as follows:

```
#define MAX_MASK_WIDTH 10  
__constant__ float M[MAX_MASK_WIDTH];
```

- This is a global variable declaration and should be outside any function in the source file.

## CONSTANT MEMORY AND CACHING

- Assume that the host code has already allocated and initialized the mask in a mask M\_h array in the host memory with Mask\_Width elements.
- The contents of the M\_h can be transferred to M in the device constant memory as follows:

```
cudaMemcpyToSymbol(M, M_h, Mask_Width*sizeof(float));
```

- Note that this is a special memory copy function that informs the CUDA runtime that the data being copied into the constant memory will not be changed during kernel execution.

```
cudaMemcpyToSymbol(dest, src, size)
```

# CONSTANT MEMORY AND CACHING

```
__global__ void convolution_1D_ba sic_kernel(float *N, float *P, int Mask_Width,
int Width) {

    int i = blockIdx.x*blockDim.x + threadIdx.x;

    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);
    for (int j = 0; j < Mask_Width; j++) {
        if (N_start_point + j >= 0 && N_start_point + j < Width) {
            Pvalue += N[N_start_point + j]*M[j];
        }
    }
    P[i] = Pvalue;
}
```



# GPU Teaching Kit

Accelerated Computing



## Module 8.2 – Parallel Computation Patterns (Stencil)






Tiled Convolution

# Objective

- To learn about tiled convolution algorithms
  - Some intricate aspects of tiling algorithms
  - Output tiles versus input tiles


# Tiling Opportunity Convolution

- Calculation of adjacent output elements involve shared input elements
  - E.g.,  $N[2]$  is used in calculation of  $P[0]$ ,  $P[1]$ ,  $P[2]$ ,  $P[3]$  and  $P[4]$  assuming a 1D convolution Mask\_Width of width 5
- We can load all the input elements required by all threads in a block into the shared memory to reduce global memory accesses




# Input Data Needs

- Assume that we want to have each block to calculate  $T$  output elements
  - $T + \text{Mask\_Width} - 1$  input elements are needed to calculate  $T$  output elements
  - $T + \text{Mask\_Width} - 1$  is usually not a multiple of  $T$ , except for small  $T$  values
  - $T$  is usually significantly larger than  $\text{Mask\_Width}$





# Definition – output tile




Each thread block calculates an output tile

Each output tile width is O\_TILE\_WIDTH

For each thread,

O\_TILE\_WIDTH is 4 in this example

# Definition - Input Tiles




**Each input tile has all values needed to calculate the corresponding output tile.**

# Two Design Options

- Design 1: The size of each thread block matches the size of an output tile
  - All threads participate in calculating output elements
  - `blockDim.x` would be 4 in our example
  - Some threads need to load more than one input element into the shared memory
- Design 2: The size of each thread block matches the size of an input tile
  - Some threads will not participate in calculating output elements
  - `blockDim.x` would be 8 in our example
  - Each thread loads one input element into the shared memory
- We will present Design 2 and leave Design 1 as an exercise.

# Thread to Input and Output Data Mapping



For each thread,  
 $\text{Index}_i = \text{index}_o - n$

were  $n$  is  $\text{Mask\_Width} / 2$   
 $n$  is 2 in this example

# All Threads Participate in Loading Input Tiles

```
float output = 0.0f;

if((index_i >= 0) && (index_i < Width)) {
    Ns[tx] = N[index_i];
}
else{
    Ns[tx] = 0.0f;
}
```

## Some threads do not participate in calculating output

```
if (threadIdx.x < O_TILE_WIDTH) {  
    output = 0.0f;  
    for(j = 0; j < Mask_Width; j++) {  
        output += M[j] * Ns[j+threadIdx.x];  
    }  
    P[index_o] = output;  
}
```

- $\text{index\_o} = \text{blockIdx.x} \times \text{O\_TILE\_WIDTH} + \text{threadIdx.x}$
- Only Threads 0 through  $\text{O\_TILE\_WIDTH}-1$  participate in calculation of output.

# Shared Memory Data Reuse

**N\_ds**

Mask\_Width is 5



Element 2 is used by thread 4 (1X)

Element 3 is used by threads 4, 5 (2X)

Element 4 is used by threads 4, 5, 6 (3X)

Element 5 is used by threads 4, 5, 6, 7 (4X)

Element 6 is used by threads 4, 5, 6, 7 (4X)


Element 7 is used by threads 5, 6, 7 (3X)

Element 8 is used by threads 6, 7 (2X)

Element 9 is used by thread 7 (1X)

# Ghost Cells

# N





# An 8-element Convolution Tile

**N\_ds**



**P**


Mask\_Width is 5



For Mask\_Width=5, we load  $8+5-1=12$  elements  
(12 memory loads)

# Each output P element uses 5 N elements

**N\_ds**



P[8] uses N[6], N[7], N[8], N[9], N[10]

P[9] uses N[7], N[8], N[9], N[10], N[11]

P[10] uses N[8], N[9], N[10], N[11], N[12]

...

P[14] uses N[12], N[13], N[14], N[15], N[16]

P[15] uses N[13], N[14], N[15], N[16], N[17]

# A simple way to calculate tiling benefit

- $(8+5-1)=12$  elements loaded
- $8*5$  global memory accesses replaced by shared memory accesses
- This gives a bandwidth reduction of  $40/12=3.3$

# Examples of Bandwidth Reduction for 1D

The reduction ratio is:

$$\text{MASK\_WIDTH} * (\text{O\_TILE\_WIDTH}) / (\text{O\_TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)$$

O_TILE_WIDTH	16	32	64	128	256
MASK_WIDTH= 5	4.0	4.4	4.7	4.9	4.9
MASK_WIDTH = 9	6.0	7.2	8.0	8.5	8.7



# GPU Teaching Kit

Accelerated Computing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).